

# YoungEverest — Bike Sharing Demand Updated Midpoint Report

---

## 1. Updated Dataset Description & Data Cleaning

The dataset used is the Bike Sharing Demand dataset from Kaggle. It contains 10,886 hourly observations, including:

- Weather variables (temperature, humidity, windspeed)
- Calendar indicators (season, holiday, workingday)
- Timestamp information
- Total bike rentals (count), which is our regression target

### Data Cleaning steps performed

1. **Removed leakage features** (casual and registered) because they sum directly to the target count.
2. Parsed the datetime column and extracted: hour, weekday, month, and year
3. Verified there were **no missing values**.
4. Reintroduced important categorical variables (season, weather) and encoded them using **One-Hot Encoding**.
5. Applied **StandardScaler** for linear models only.
6. Ensured **train/validation/test split** uses stratification and a fixed random seed.

## 2. Exploratory Data Analysis (EDA)


This section includes the two required midpoint plots.

### Plot 1: Target Distribution for Classification

#### Description:

A peak hour is defined as any hour where count  $\geq 100$ .

This threshold is supported by the histogram and hourly usage pattern, where nighttime hours consistently fall below 100, and daytime/commute hours frequently exceed it.



## Plot 2: Correlation Heatmap (Key Numeric Features)

### Description:

Temperature and hour show strong positive correlation with bike rentals, while humidity shows a negative correlation. These relationships motivated feature choices for baseline models.

## 3. Train–Validation–Test Split

The dataset was split into:

- 70% Training set (7,620 rows)
- 15% Validation set (1,633 rows)
- 15% Test set (1,633 rows)

A **fixed random seed (42)** ensures full reproducibility.

## 4. Baseline Models with MLflow Tracking

We implemented four classical models:

### Classification Baselines

1. Logistic Regression (scaled)
2. Decision Tree Classifier (unscaled)

## Regression Baselines

1. Linear Regression (scaled)
2. Decision Tree Regressor (unscaled)

All models were logged using **MLflow**, including: parameters, metrics, model artifacts.

## 5. Results and Discussion

**Table 1 — Classification Metrics (Validation + Test)**

Model	Accuracy (Test)	F1 (Test)
Logistic Regression	<b>0.8953</b>	<b>0.9129</b>
Decision Tree Classifier	<b>0.8959</b>	<b>0.9154</b>

### Summary:

The Decision Tree Classifier slightly outperforms Logistic Regression.

Tree-based models capture non-linear relationships such as:

- morning and evening commute peaks
- weather effects
- sudden changes in demand during extreme humidity

**Table 2 — Regression Metrics (Validation + Test)**

Model	MAE (Test)	RMSE (Test)
Linear Regression	<b>82.70</b>	<b>111.31</b>
Decision Tree Regressor	<b>51.52</b>	<b>84.63</b>

### Summary:


The Decision Tree Regressor significantly outperforms Linear Regression.

Linear Regression struggles with non-linearity, whereas trees adapt to sharp demand changes across hours and seasons.

## 6. Required Plots for Model Evaluation

### Plot 3 — Confusion Matrix (Best Classifier: Decision Tree)

Shows strong classification performance with balanced predictions across peak and off-peak classes.



### Plot 4 — Residuals vs Predicted (Best Regressor: Decision Tree)

Residuals are close to zero for moderate demand and show expected variance for extreme values.

## 7. Discussion: What Worked & Failure Modes

### What Worked

- Decision trees captured complex interactions (weather  $\times$  hour  $\times$  temperature).
- One-hot encoding fixed earlier categorical processing issues.
- Scaling improved convergence for linear/logistic models.
- Stratified splits maintained class balance and improved stability.
- MLflow ensured complete reproducibility and proper experimental tracking.

### Failure Modes / Limitations

- Linear models struggled with non-linear patterns and high variance.
- Residuals increase at higher rental counts (heteroscedasticity).
- Decision trees can overfit; deeper trees showed unstable validation metrics.

## 8. Neural Network Plan (For Final Report)

We will implement a **Multilayer Perceptron (MLP)** for both regression and classification:

- 2–3 hidden layers (64–128 neurons each)
- ReLU activation
- Batch normalization
- Dropout for regularization
- Output layers:
  - Sigmoid for classification
  - Linear for regression
- Adam optimizer + early stopping

### Justification:

MLPs handle non-linear relationships better than linear models and reduce the overfitting tendency of single decision trees.

## 9. Appendix — MLflow Experiment Tracking

This screenshot confirms that all baseline models were successfully logged with parameters, metrics, and artifacts, meeting the reproducibility requirement.

The screenshot shows the MLflow interface with the following details:

- Header:** mlflow 3.5.1, GitHub, Docs
- Left Sidebar:** Home, Experiments (selected), Models, Prompts
- Current Experiment:** BikeSharing\_midpoint (Machine learning)
- Table Headers:** Run Name, Created, Dataset, Duration, Source, Models
- Data Rows:** 16 rows of run information, each with a checkbox, a colored circle, and a truncated run name. Most runs are labeled "decision\_tree\_regressor" or "linear\_regression\_bas...".
- Table Footer:** Show more columns (10 total)
- Bottom:** 16 matching runs

## **10. References**

1. Fanaee-T, H., & Gama, J. (2014). Event labeling combining ensemble detectors and background knowledge. *Progress in Artificial Intelligence*, 2(2–3), 113–127.
2. Pedregosa, F., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
3. MLflow Documentation. <https://mlflow.org/>
4. Hunter, J. D. (2007). Matplotlib: A 2D Graphics Environment. *Computing in Science & Engineering*, 9(3), 90–95.
5. Waskom, M. L. (2021). Seaborn: Statistical Data Visualization. *Journal of Open Source Software*, 6(60), 3021.

**GitHub Repository - <https://github.com/arjun-sapkota887/Bike-Sharing-Demand>**